

Part II: Reinforcement learning

Challenges of reinforcement learning (recap from slides01)

Reinforcement learning is the subfield of machine learning that deals with learning from **interaction** with an unknown environment

Unlike supervised learning:

- the agent only gets feedback on the *chosen* action;
- the feedback is delayed: it applies to sequences of actions;
- by choosing actions, the agent is in control of what data is seen.



Exploration–Exploitation Dilemma

At any point, the agent will have limited information, and can choose actions to:

- **exploit** that information: pick the action believed to be best
- **explore**: try out another action, maybe getting new information

In the long term, the best performance will be obtained if there is a good balance between exploring and exploiting

Multi-armed bandits: introduction

Bandit problems are a simplified type of reinforcement problem:

- The environment has only one *state*
- but the exploration–exploitation dilemma is there (all the more clearly)

The simplification allows these problems to be thoroughly analysed mathematically, though many important questions are still unanswered

- The part of chapter 2 we will skip (starting with section 2.5) goes into some of this theory

Multi-armed bandits: definition

(Named after slot machines: “one-armed bandits”)

Each time step t , the agent chooses an action A_t and receives reward R_t

- k available actions
- $q_*(a) := \mathbb{E}[R_t | A_t = a]$: expected reward of action a (also called action **value**)
(*’s will denote ‘true’ or ‘optimal’)




Taking the action with the largest $q_*(a)$ would give the optimal reward... but we don't know those values

- $Q_t(a)$: our estimate of $q_*(a)$ at time t (before taking an action)
 - Obvious choice: the average of the R_i 's at all $i < t$ where $A_i = a$

Taking the action with the largest $Q_t(a)$ is called the **greedy** action

Example



Adding exploration

By taking the greedy action, we are doing only exploitation, no exploration

Straightforward way to ensure we do some exploration as well: for each action, with some small probability ϵ , select an action uniformly at random

- Note that this will still end up selecting the greedy action, with overall probability ϵ/k

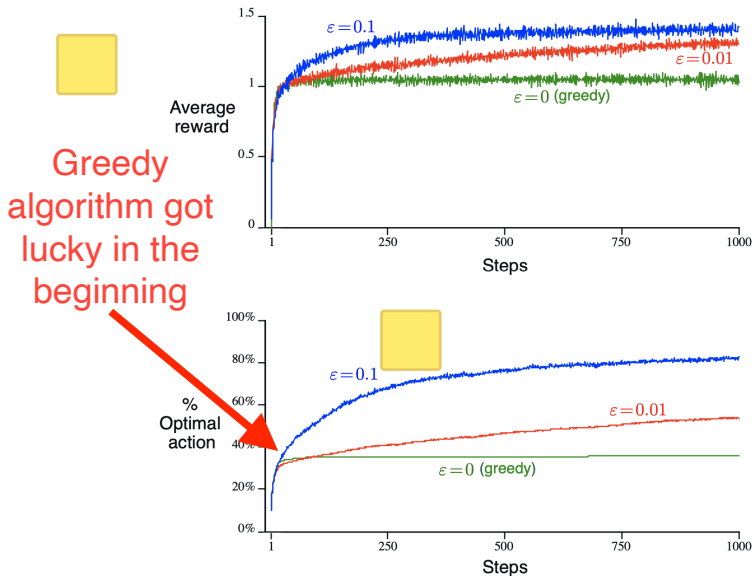
This is called an ϵ -greedy strategy



Why does this work?

- As t goes to infinity, so does the number of times each action is tried. So our estimates $Q_t(a)$ will converge to $q_*(a)$.

Example results (over 2000 different problems)



Observations

- ϵ too small (e.g. 0): likely to get stuck with suboptimal action
- ϵ too large: in the long run, keeps making wrong moves even if it has already found the optimal action

Good value depends very much on the problem. E.g. if R_t has small variance, it will be easier to identify the optimal action, so a small ϵ is good enough



Improvements:

- We could let ϵ get smaller over time
- Many more ways to improve over this simple idea (e.g. by taking into account how confident we are about each estimate, or how likely some action is to turn out optimal)



Incrementally tracking value estimates

Suppose R_1, \dots, R_{n-1} are previous rewards obtained after selecting action a (note the simplified indexing on this slide). Then


$$Q_n := \frac{R_1 + \dots + R_{n-1}}{n-1}$$

No need to keep track of all numbers R_1, \dots, R_{n-1} . Instead, update Q_n to Q_{n+1} when new reward R_n comes in:

$$\begin{aligned} Q_{n+1} &= \frac{1}{n} \sum_{i=1}^n R_i \\ &= \frac{1}{n} \left(R_n + (n-1) \frac{1}{n-1} \sum_{i=1}^{n-1} R_i \right) \\ &= \frac{1}{n} (R_n + (n-1)Q_n) \\ &= Q_n + \frac{1}{n} [R_n - Q_n] \end{aligned}$$

Incremental formula

$$Q_{n+1} = Q_n + \frac{1}{n} [R_n - Q_n]$$

- Cheap in terms of memory and computation 
- Example of a general form we will see very often:

NewEstimate = OldEstimate + StepSize(Target - OldEstimate)



Simple bandit algorithm



A simple bandit algorithm

Initialize, for $a = 1$ to k :

$$Q(a) \leftarrow 0$$

$$N(a) \leftarrow 0$$

Loop forever:

$$A \leftarrow \begin{cases} \operatorname{argmax}_a Q(a) & \text{with probability } 1 - \varepsilon \quad (\text{breaking ties randomly}) \\ \text{a random action} & \text{with probability } \varepsilon \end{cases}$$


$$R \leftarrow \text{bandit}(A)$$

$$N(A) \leftarrow N(A) + 1$$



$$Q(A) \leftarrow Q(A) + \frac{1}{N(A)} [R - Q(A)]$$




Markov Decision Process

The general reinforcement learning problem can be formalized by a Markov decision process (MDP) 

Compared to bandit problems: the environment can now be in one of multiple states s



-  Could reflect raw sensor readings
-  Could also reflect other information the agent has about the environment

Has been studied extensively, e.g. in mathematics and engineering (where it also called *control*) 


MDPs: additional challenge

- The agent's action may influence the state of the environment
- So we may need to choose between large immediate reward but going to a less desirable state, or small immediate reward but a better chance of going to a better state in the future

Reward, state, action



Note that the index t increases *after* the action

- $\mathcal{A}(s)$: set of actions available in state s 
 - Even if 'the same' action is available in different states, it might be good in one but bad in another state. So we can think of them as being different actions that we need to learn about separately.
- Dynamics function (suggestively written with a $|$):

$$p(s', r | s, a) := P(S_t = s', R_t = r | S_{t-1} = s, A_{t-1} = a)$$

'Markov' decision process: new state and reward determined by preceding state and action alone (and in particular not by anything longer ago)

Example: recycling robot

from the book

s	a	s'	$p(s' s, a)$	$r(s, a, s')$
high	search	high	α	r_{search}
high	search	low	$1 - \alpha$	r_{search}
low	search	high	$1 - \beta$	-3
low	search	low	β	r_{search}
high	wait	high	1	r_{wait}
high	wait	low	0	-
low	wait	high	0	-
low	wait	low	1	r_{wait}
low	recharge	high	1	0
low	recharge	low	0	-

