

Reward hypothesis: “All of what we mean by goals and purposes can be well thought of as the *maximization* of the *expected value* of the *cumulative sum* of a received scalar signal (called reward)”

So in an MDP, the reward should represent the *goal* we want the agent to accomplish


- Not e.g. some information about *how* it should reach the goal



Episodic tasks

Suppose the interaction ends after a finite number of time steps (e.g. in tasks where a goal is eventually reached)

- Each such subsequence of time steps is called an **episode**
- Termination is marked by transitioning to a special state: the **terminal state**; all transitions from this state go back to this state, and give reward 0
- State set notation: $\mathcal{S}^+ := \mathcal{S} \cup \{\text{terminal state}\}$



After an episode is terminated, we can think of the agent as starting a new episode (where it can apply what it learned in previous episodes)

Returns in episodic tasks



Return: rewards from time t to the end

$$G_t := R_{t+1} + R_{t+2} + \dots + R_T$$

Here T is the termination time (a random variable): the time step where the terminal state is first reached

So:

- reward: in one time step
- return: until the end
- value: expected return

Continuing tasks

If a task might continue infinitely long, it is called a **continuing task**

Issue: sum of rewards might be infinite/undefined!

Solution: use **discounting** with some chosen rate $0 \leq \gamma \leq 1$, and define G_t as

$$G_t := R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

Expresses that rewards farther in the future should be given less weight

Special cases:

- $\gamma = 1$: no discounting, ordinary sum
- $\gamma = 0$: care only about immediate future (take $0^0 = 1$; called “myopic”)

A **policy** specifies how an agent chooses actions

- May be probabilistic in general (like the ϵ -greedy strategies we saw before)
- Notation: $\pi(a | s) := P(A_t = a | S_t = s)$ (for all t)

If we know p and π , we can compute probabilities of all future events given a current state

Exercise 3.11 (section 3.5): If the current state is S_t , and actions are selected according to stochastic policy π , then what is the expectation of R_{t+1} in terms of π and the four-argument function p (3.2)? ($p(s', r | s, a)$)

Value functions

For known p and π , the **value** (expected return) is described for each current state by



$$v_{\pi}(s) := E_{\pi}[G_t \mid S_t = s]$$

(the **state-value function for policy π**)

Similarly, if both S_t and A_t are known we have the **action-value function for policy π** :



$$q_{\pi}(s, a) := E_{\pi}[G_t \mid S_t = s, A_t = a]$$


Bellman equations

Rewriting $v_\pi(s)$, we find

$$\begin{aligned} v_\pi(s) &:= E_\pi[G_t \mid S_t = s] \\ &= E_\pi[R_{t+1} + \gamma G_{t+1} \mid S_t = s] \\ &= \sum_a \pi(a \mid s) \sum_{s', r} p(s', r \mid s, a) [r + \gamma E_\pi[G_{t+1} \mid S_{t+1} = s']] \\ &= \sum_a \pi(a \mid s) \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_\pi(s')] \end{aligned}$$

Called the **Bellman equation for v_π**

Plays a central role in many reinforcement learning methods



Optimal policies and value functions

A policy π is at least as good as another policy π' if

$$v_{\pi}(s) \geq v_{\pi'}(s) \quad \text{for all } s$$

A policy that is at least as good as *all* other policies is an **optimal policy** (written π_*)

- one always exists!
- there may be more than one

All optimal policies have the same **optimal (state/action-)value functions**:



$$v_*(s) := \max_{\pi} v_{\pi}(s)$$



$$q_*(s, a) := \max_{\pi} q_{\pi}(s, a)$$


Bellman optimality equation

Bellman equation for v_* :

$$\begin{aligned}v_*(s) &= \max_{a \in \mathcal{A}(s)} q_{\pi_*}(s, a) \\&= \max_a E_{\pi_*}[G_t \mid S_t = s, A_t = a] \\&= \max_a E_{\pi_*}[R_{t+1} + \gamma G_{t+1} \mid S_t = s, A_t = a] \\&= \max_a E[R_{t+1} + \gamma v_*(S_{t+1}) \mid S_t = s, A_t = a] \\&= \max_a \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_*(s')]\end{aligned}$$

(and similar for q_*)

In backup diagrams,
arcs represent taking
maximum:



Computing optimal policies and value functions

- Using the Bellman equations, it is in principle possible to compute v_* and q_*
 - Assuming p is known
 - May take a prohibitive amount of computation time and memory if number of states / actions is large
- Once we know v_* , we can compute an optimal policy: for each state, compare each action's expected immediate reward plus the expected return after that (i.e. greedy w.r.t. v_*)
- If we know q_* , this becomes even easier



In reinforcement learning, p is rarely known in practice. Approach: estimate it from experience, and use it to in turn approximate $v_*/q_*/\pi_*$

Prediction



prediction: problem of finding an estimate V of v_π for a given policy π

Possible approaches:

- Dynamic programming (chapter 4; assumes p is known) 
 - Keep track of our estimates V so far
 - keep improving (for all states s) our estimates $V(s)$ by using new estimates $V(s')$ for successor states
- Monte Carlo methods (chapter 5) 
 - Play through an episode until termination
 - for each state s visited during the episode, we know the return we got, and can use this to improve our estimate $V(s)$

Prediction in temporal-difference learning



Temporal-difference (TD) learning: combines ideas from dynamic programming (DP) and Monte Carlo (MC) methods:

- Like MC, based on experience rather than knowledge of p
- Unlike MC, don't wait for an episode to end before we update V

One-step TD

Tabular TD(0) for estimating v_π

Input: the policy π to be evaluated

Algorithm parameter: step size $\alpha \in (0, 1]$

Initialize $V(s)$, for all $s \in \mathcal{S}^+$, arbitrarily except that $V(\text{terminal}) = 0$

Loop for each episode:

Initialize S

Loop for each step of episode:

$A \leftarrow$ action given by π for S

Take action A , observe R, S'

$V(S) \leftarrow V(S) + \alpha [R + \gamma V(S') - V(S)]$

$S \leftarrow S'$

until S is terminal

estimate on return from
now on until the end



Notes:

- $V(S)$ is updated using a version of the incremental formula from chapter 2
- Like DP, TD *bootstraps*: uses previous estimates for other states to form a new estimate (MC doesn't do this)

Control: generalized policy iteration

Next to *prediction*, there is the **control** problem: approximating optimal policies



Common idea for DP/MC/TD: **generalized policy iteration (GPI)**

- keep track of a policy π and estimated (action) values Q
- repeatedly improve the estimate Q for the current policy π (the prediction problem)
- at the same time, repeatedly improve π by making it closer to the greedy strategy w.r.t. Q

TD control: Sarsa

Estimates Q instead of V , using a similar formula:

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

Named **Sarsa** because it uses $(S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1})$

Sarsa algorithm

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Loop for each step of episode:

 Take action A , observe R , S'


 Choose A' from S' using policy derived from Q (e.g., ε -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'$; $A \leftarrow A'$;

 until S is terminal

On- vs. off-policy control

Sarsa estimates Q for the policy it's currently following: example of an **on-policy** method 

Sometimes we want to estimate values for a different policy than the one we're following

- Common reason: the policy we're following needs to do exploration, but we know this makes it worse at exploitation

Q-learning

Q-learning: estimate the optimal policy's q_* , regardless of the policy being followed

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+$, $a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R , S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

 until S is terminal

Expected Sarsa



Expected Sarsa: estimate any **target policy**'s q_π , regardless of the policy being followed

Update rule:

$$\begin{aligned} &Q(S_t, A_t) \\ &\leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma E_\pi [Q(S_{t+1}, A_{t+1}) | S_{t+1}] - Q(S_t, A_t)] \end{aligned}$$

More general than Q-learning